


Online Resource 3



Caption. Debate-level accuracy of three classical models (weighted sum, linear discriminant, decision forest) reported in Table 9 of the main text, evaluated two ways. Holding out the debate leaves both of its speakers in the training data; holding out the speakers removes them entirely. All three models fall to chance under the stricter test. This figure presents, in visual form, the same comparison tabulated in Table 9 of the main text.