

Online Resource 4

Input	Outcome accuracy	Exact debate identified
Descriptors only	13/24	0/24
Transcript only	24/24	24/24

Caption. The language model’s debate-level accuracy (majority vote over 3 re-randomised repeats) beside how often a separate, question-free probe identifies the actual debate from the same input. These are the same two numbers plotted in Figure 4 of the main text (“Transcript-Only and Descriptor-Only Judgments” / “Contamination Probe”).